

problemset1a.R

anirudh

2023-09-19

```
library(tidyverse)

## -- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
## v dplyr      1.1.2      v readr      2.1.4
## v forcats    1.0.0      v stringr   1.5.0
## v ggplot2    3.4.3      v tibble    3.2.1
## v lubridate  1.9.2      v tidyr     1.3.0
## v purrr      1.0.2
## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()     masks stats::lag()
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors

# First steps - R
# Introduction to R & Data Visualization

# Part 1: R Questions
# Question 1: First Steps with R
# Part 4: Generating sequences of numbers and logical vectors
# 1.
seq(1:15)

## [1]  1  2  3  4  5  6  7  8  9 10 11 12 13 14 15

# 2.
2*1:15

## [1]  2  4  6  8 10 12 14 16 18 20 22 24 26 28 30

(2*1):15

## [1]  2  3  4  5  6  7  8  9 10 11 12 13 14 15

# There is a difference here. In the first one, it's skipping the value by 2 numbers.
# In the second one, it's starting from the number 2 instead of 1.
# 3.
# A logical vector only consists of values 0 and 1 or True or False
vec <- runif(10, min = 0, max = 3)
vec_more_1.5 <- vec[vec > 1.5]
vec_more_1.5

## [1] 2.403222 2.209297 2.622650 2.898517

# 4.
vec_equal_1.2 <- vec[vec == 1.2]
vec_equal_1.2
```

```
## numeric(0)
# I was expecting a 0 because 1.2 is a very specific number to expect in a randomly generated vector
# 5.
?seq
seq(from = 15, to = 1)

## [1] 15 14 13 12 11 10 9 8 7 6 5 4 3 2 1

# Part 5: Arrays & Matrices
# 1.
x <- runif(100)
x

## [1] 0.659478828 0.380226362 0.800692990 0.785761222 0.600894168 0.904707217
## [7] 0.284712873 0.725831782 0.318796805 0.851004535 0.807857111 0.814198268
## [13] 0.734832927 0.283199710 0.924963735 0.162544179 0.342285622 0.990165519
## [19] 0.702170232 0.630128546 0.253773024 0.308152263 0.043414640 0.227923824
## [25] 0.973692215 0.019666349 0.007388633 0.210812387 0.019205903 0.583211335
## [31] 0.913057485 0.174814483 0.052148457 0.924308672 0.938043522 0.833015616
## [37] 0.830134898 0.326950809 0.260545353 0.070827676 0.190997397 0.342037965
## [43] 0.879228737 0.624236160 0.316455274 0.210685254 0.164217372 0.462215300
## [49] 0.310923022 0.576316157 0.238342189 0.848432293 0.329860216 0.690994091
## [55] 0.058853804 0.047236713 0.083546347 0.115713447 0.327155857 0.024377362
## [61] 0.975626413 0.518508662 0.050742500 0.623748417 0.016563896 0.893278242
## [67] 0.228000259 0.265731229 0.578588570 0.842822355 0.104046014 0.406164668
## [73] 0.594460348 0.254221747 0.207385767 0.289859073 0.501668638 0.943765767
## [79] 0.296401632 0.697809929 0.635466391 0.880022825 0.326436355 0.328959376
## [85] 0.823319244 0.373875335 0.579558837 0.194886493 0.800521270 0.343988652
## [91] 0.576767597 0.314426817 0.864148528 0.322703713 0.658320538 0.663131212
## [97] 0.735854509 0.203508660 0.970714187 0.823556699

# 2.
matrix_x <- matrix(x, nrow = 10)
matrix_x

##           [,1]      [,2]      [,3]      [,4]      [,5]      [,6]      [,7]
## [1,] 0.6594788 0.8078571 0.253773024 0.91305749 0.1909974 0.23834219 0.9756264
## [2,] 0.3802264 0.8141983 0.308152263 0.17481448 0.3420380 0.84843229 0.5185087
## [3,] 0.8006930 0.7348329 0.043414640 0.05214846 0.8792287 0.32986022 0.0507425
## [4,] 0.7857612 0.2831997 0.227923824 0.92430867 0.6242362 0.69099409 0.6237484
## [5,] 0.6008942 0.9249637 0.973692215 0.93804352 0.3164553 0.05885380 0.0165639
## [6,] 0.9047072 0.1625442 0.019666349 0.83301562 0.2106853 0.04723671 0.8932782
## [7,] 0.2847129 0.3422856 0.007388633 0.83013490 0.1642174 0.08354635 0.2280003
## [8,] 0.7258318 0.9901655 0.210812387 0.32695081 0.4622153 0.11571345 0.2657312
## [9,] 0.3187968 0.7021702 0.019205903 0.26054535 0.3109230 0.32715586 0.5785886
## [10,] 0.8510045 0.6301285 0.583211335 0.07082768 0.5763162 0.02437736 0.8428224
##           [,8]      [,9]      [,10]
## [1,] 0.1040460 0.6354664 0.5767676
## [2,] 0.4061647 0.8800228 0.3144268
## [3,] 0.5944603 0.3264364 0.8641485
## [4,] 0.2542217 0.3289594 0.3227037
## [5,] 0.2073858 0.8233192 0.6583205
## [6,] 0.2898591 0.3738753 0.6631312
## [7,] 0.5016686 0.5795588 0.7358545
## [8,] 0.9437658 0.1948865 0.2035087
## [9,] 0.2964016 0.8005213 0.9707142
```

```
## [10,] 0.6978099 0.3439887 0.8235567
dim(matrix_x)

## [1] 10 10

# 3.
y <- matrix(runif(100), 10, 10)
# 4.
# entire column
y[0,]

##      [,1] [,2] [,3] [,4] [,5] [,6] [,7] [,8] [,9] [,10]
# entire row
y[,0]

##
## [1,]
## [2,]
## [3,]
## [4,]
## [5,]
## [6,]
## [7,]
## [8,]
## [9,]
## [10,]

# single element
y[0,1]

## numeric(0)
# range of rows or columns
y[0:2, 0:4]

##      [,1]      [,2]      [,3]      [,4]
## [1,] 0.7968570 0.9212767 0.6823411 0.4418082
## [2,] 0.1728649 0.2200920 0.2215528 0.6061150

# 5.
t(y)

##      [,1]      [,2]      [,3]      [,4]      [,5]      [,6]      [,7]
## [1,] 0.79685701 0.1728649 0.87578904 0.6959279 0.7600220 0.23532950 0.77857844
## [2,] 0.92127670 0.2200920 0.91410048 0.2217077 0.5799692 0.89248025 0.70558761
## [3,] 0.68234114 0.2215528 0.87926976 0.4350800 0.8786261 0.74051907 0.76129400
## [4,] 0.44180819 0.6061150 0.92364682 0.9502615 0.5233822 0.29902525 0.50676110
## [5,] 0.09774047 0.9485971 0.32964212 0.4493628 0.1469271 0.68934373 0.61363313
## [6,] 0.57022465 0.5548237 0.98581614 0.3519196 0.9349678 0.06831423 0.09488395
## [7,] 0.79057356 0.2556447 0.09482889 0.9727837 0.9473751 0.19442846 0.22069528
## [8,] 0.78001966 0.3250503 0.50257397 0.9632758 0.4546930 0.26004971 0.16289764
## [9,] 0.17029558 0.4724135 0.31265122 0.8593613 0.5559640 0.23818317 0.72675383
## [10,] 0.01520497 0.1229970 0.85465740 0.7034005 0.6937164 0.63232937 0.50927153
##      [,8]      [,9]      [,10]
## [1,] 0.81813397 0.4586118 0.82407474
## [2,] 0.86122413 0.8053837 0.33005530
## [3,] 0.50251328 0.8392831 0.04591512
```

```
## [4,] 0.56156307 0.2405766 0.06305951
## [5,] 0.11714718 0.8932037 0.76804530
## [6,] 0.34385035 0.9930892 0.39536212
## [7,] 0.40769378 0.3726629 0.38666612
## [8,] 0.06807029 0.3033542 0.68106720
## [9,] 0.23121343 0.2542687 0.39967714
## [10,] 0.34928806 0.9481925 0.95433873

# 6.
a <- matrix(runif(4), 2, 2)
b <- matrix(runif(4), 2, 2)
a*b # element wise

##           [,1]      [,2]
## [1,] 0.1719298 0.2987201
## [2,] 0.1151530 0.6396910

a %*% b # mathematical matrix multiplication

##           [,1]      [,2]
## [1,] 0.5444728 0.4937583
## [2,] 0.8017800 0.7320253

# 7.
rm(list=ls())

# Part 6: The Working Directory and R Projects
getwd()


## [1] "/home/anirudh/University/E401 Machine Learning For Economic Data/Week 3/problemset1a"
setwd("/home/anirudh/University/E401 Machine Learning For Economic Data/Week 3/problemset1a/r_firststep")
vgr = read_csv("vgr.csv")

## Rows: 140 Columns: 4
## -- Column specification -----
## Delimiter: ","
## dbl (4): con, p, yv, quarter
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.


# Problem Set - 1a

# Question 2: Data Visualization
mpg <- mpg
diamonds <- diamonds

# Question 2.1: Visualization Basics
ggplot(data = mpg)
```



```
# I see a gray box here with no plot being displayed because we aren't providing the structure  
# information of the plot to the ggplot command  
?mpg  
# The drv variable is short for drive which describes the type of drive mode the car uses. For example,  
# front-wheel, drive, rear wheel drive, or 4 = 4wd  
ggplot(data = mpg) + geom_point(mapping = aes(x = drv, y = class))
```




We see a plot categorizing vehicles into specific categories like SUV, pickup truck based on their drive type. However, due to overlapping points, we don't see all 234 vehicle entries, making it challenging to get a clear understanding of how the actual number of vehicles are categorized.


Question 2.2: Aesthetic Mappings

1.


```
ggplot(data = mpg) + geom_point(mapping = aes(x = displ, y = hwy, color = "blue"))
```



```
# The reason it doesn't show blue colour is because the "color" argument is not part of the command aes
# The below line of code shows how it should actually be typed for the colour blue to be used.
ggplot(data = mpg) + geom_point(mapping = aes(x = displ, y = hwy), color = "blue")
```




```
# 2.  
?mpg  
# categorical variables: manufacturer, model, drv, class, trans, fl  
# continuous variables: displ, year, cyl, cty, hwy  
  
# 3.  
ggplot(mpg, aes(x = displ, y = hwy, colour = cty)) + geom_point()
```

cty here is a continuous variable that colours the points from light to dark. Light for points high and dark for points low.


```
ggplot(mpg, aes(x = displ, y = hwy, size = cty)) + geom_point()
```



when cty is used for size, the size values become continuous according to the corresponding y value's

4.

```
ggplot(mpg, aes(x = displ, y = hwy, colour = hwy, size = displ)) + geom_point()
```




*# The issue with mapping a single variable to multiple aesthetics is that there is redundancy. So it's
avoid this approach.*

Question 2.3: Facets

#1.


```
ggplot(mpg, aes(x = displ, y = hwy)) + geom_point() + facet_grid(. ~ cty)
```




If we facet a continuous variable, we can see that it gets converted into a categorical variable.

#2.

```
ggplot(data = mpg) + geom_point(mapping = aes(x = drv, y = cyl))
```

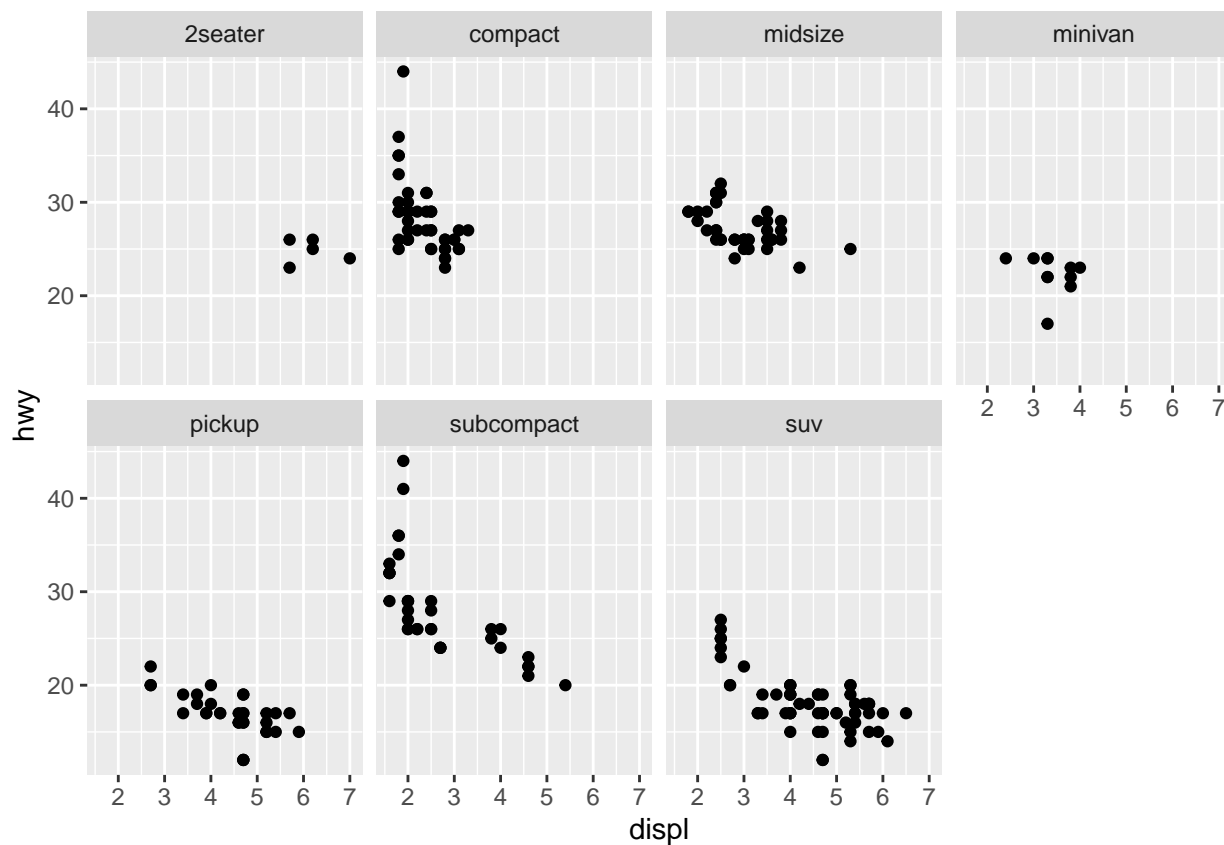


```
ggplot(data = mpg) + geom_point(mapping = aes(x = drv, y = cyl)) + facet_grid(drv ~ cyl)
```



*# The empty cells in this plot are data points in drv vs cyl that have no observations.
These are the same locations in the scatter plot of drv and cyl that have no points.*

```
ggplot(data = mpg) + geom_point(mapping = aes(x = displ, y = hwy)) + facet_wrap(~ class, nrow = 2)
```



The advantage would be we can encode them into distinct categorical variables
The disadvantage would be that it makes it difficult to compare the datapoints between categories
But as the no. of categories increase, differences in colours decrease making it difficult to tell the
datapoints apart

Question 2.4:

1.

For line chart, geom_line()

For boxplot, geom_boxplot()


For histogram, geom_histogram()

For area chart, geom_area()


#2.

```
ggplot(data = mpg, mapping = aes(x = displ, y = hwy)) +  
  geom_point() +  
  geom_smooth()
```

```
## `geom_smooth()` using method = 'loess' and formula = 'y ~ x'
```



```
ggplot() +  
  geom_point(data = mpg, mapping = aes(x = displ, y = hwy)) +  
  geom_smooth(data = mpg, mapping = aes(x = displ, y = hwy))  
  
## `geom_smooth()` using method = 'loess' and formula = 'y ~ x'
```

They don't look different because geom_point and smooth because the dataset is still the same

Part 2: Your Project

Question 3:

`getwd()`


```
## [1] "/home/anirudh/University/E401 Machine Learning For Economic Data/Week 3/problemset1a/r_firststep"
```

```
setwd("/home/anirudh/University/E401 Machine Learning For Economic Data/Week 3/problemset1a/part_2_data")
```

```
spotify <- read.csv("playlist_2010to2022.csv")
```

```
ggplot(data = spotify) + geom_point(mapping = aes(x = duration_ms, y = track_popularity))
```

```
## Warning: Removed 1 rows containing missing values (`geom_point()`).
```




*# The no.of milliseconds seems to play a role in the songs popularity. It looks like songs which have a
around 200000ms/200s/3.33 minutes are popular.*

```
musical_attributes <- c("danceability", "energy", "key", "loudness", "mode", "speechiness", "acousticness", "instrumentalness", "liveness", "valence", "tempo")
musical_attributes
```

```
## [1] "danceability" "energy" "key" "loudness"
## [5] "mode" "speechiness" "acousticness" "instrumentalness"
## [9] "liveness" "valence" "tempo"
```


```
ggplot(data = spotify) + geom_point(mapping = aes(x = track_popularity, y = danceability))
```

```
## Warning: Removed 1 rows containing missing values (`geom_point()`).
```




```
ggplot(data = spotify) + geom_point(mapping = aes(x = track_popularity, y = loudness))
```

```
## Warning: Removed 1 rows containing missing values (`geom_point()`).
```




```
ggplot(data = spotify) + geom_point(mapping = aes(x = track_popularity, y = speechiness))
```

```
## Warning: Removed 1 rows containing missing values (`geom_point()`).
```




```
ggplot(data = spotify) + geom_point(mapping = aes(x = track_popularity, y = instrumentalness))
```

```
## Warning: Removed 1 rows containing missing values (`geom_point()`).
```




```
getwd()
```


```
## [1] "/home/anirudh/University/E401 Machine Learning For Economic Data/Week 3/problemset1a/part_2_data"
setwd("/home/anirudh/University/E401 Machine Learning For Economic Data/Week 3/problemset1a/part_2_data")
house <- read.csv("house.csv")
ggplot(data = house) + geom_point(mapping = aes(x = metro_distance, y = price))
```



```
ggplot(data = house) + geom_point(mapping = aes(x = bedroom_count, y = price))
```



```
ggplot(data = house) + geom_point(mapping = aes(x = age, y = price))
```

```
ggplot(data = house) + geom_point(mapping = aes(x = floor, y = price))
```

